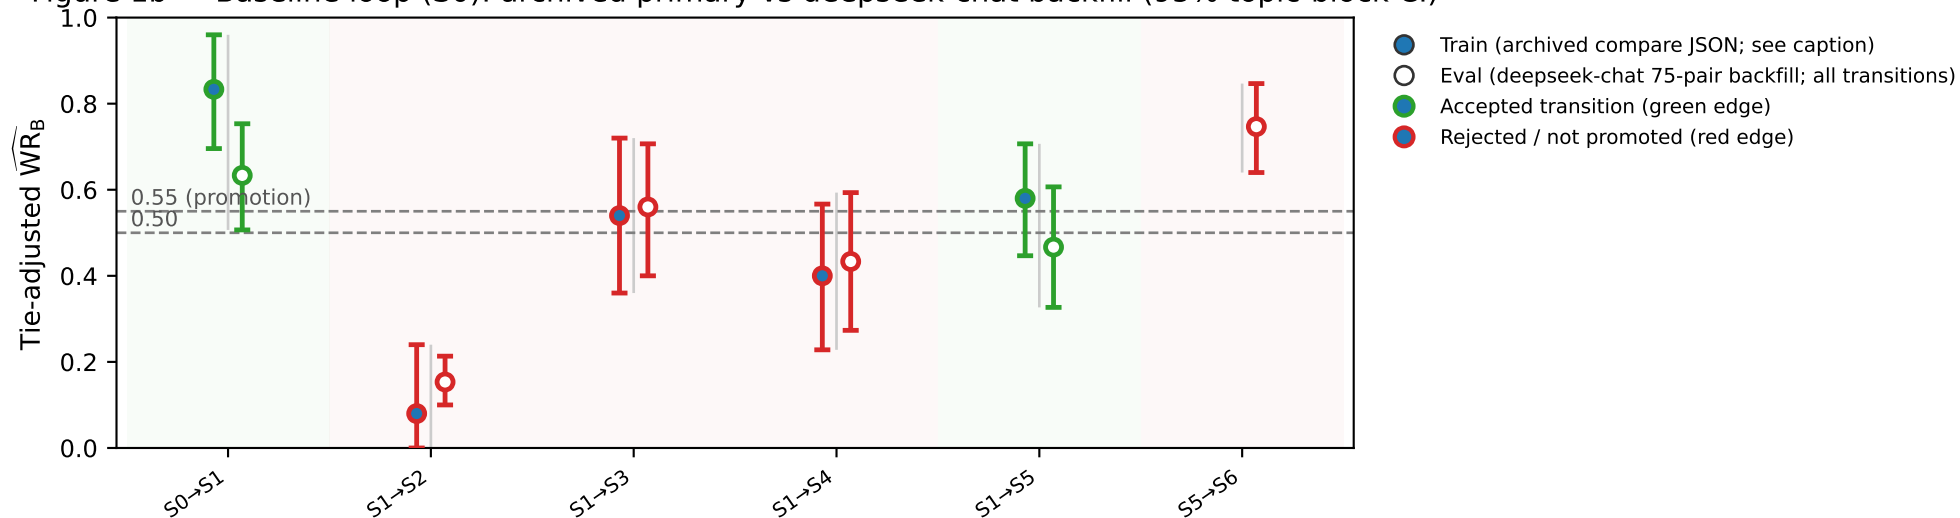


Figure 1b — Baseline loop (S0): archived primary vs deepseek-chat backfill (95% topic-block CI)



****Train (filled):**** ``compare\_<A>\_vs\_<B>.json`` restored from git (parent of ``620151b``). Pair counts and ``judge\_model`` match those historical sessions. ****S5→S6**** has no archived compare — train marker omitted (loop rebooted mid-run). ****Eval (open markers):**** full 75-pair ``blind\*\_deepseek-chat.json`` from ``ideas/blind\_backfill.py`` (same DeepSeek judge for every transition, including S5→S6). ****Note:**** backfill re-judges text from on-disk ``results/<S#>/ideas.json`` (later regeneration), so eval can diverge from the archived primary session. Bands: accept (green) / reject (red).