


S15 vs S_sota: batch win rates and judge-human alignment

Publish-eval intersection (N = 64)



Inter-judge agreement (publish-eval)

Low pair-level κ ; left-panel batch means still not near 0.5.

Pair (N = 64)	Cohen κ	Exact label match
DeepSeek ↔ GPT-5.4	0.16	53%
DeepSeek ↔ G3 Flash	0.07	41%
GPT-5.4 ↔ G3 Flash	0.21	56%

vs. human experts (custom 12-topic pilot)

Automated judge	Pearson r	Side match
DeepSeek	-0.06	42%
GPT-5.4	+0.04	25%
Gemini 3 Flash	-0.07	17%