



Intra-judge test-retest (publish-eval intersection, N = 64 pairs)

Exact label match
(same pass vs rerun)



Pairwise agreement
(beyond chance)



Second passes: ``judge\_gpt-5.4\_rerun.json`` and ``judge\_deepseek-chat\_rerun.json`` vs primaries on the same idea pairs as ``judge\_analysis.json``. Dotted line: 1/3 (uniform random 3-class match). Tie-adjusted $\widehat{WR}_B$ for `S\_sota` on this slice: GPT-5.4 0.234 $\rightarrow$ 0.227 (0.8 pp shift); DeepSeek 0.320 $\rightarrow$ 0.375 (5.5 pp shift).